


## Linux kernel map



**What are benefits  
of knowing Linux  
kernel?**



# CVE Details

The ultimate security vulnerability datasource

[Log In](#) [Register](#) [Reset Password](#) [Activate Account](#)

[Home](#)

**Browse :**

[Vendors](#)

[Products](#)

[Vulnerabilities By Date](#)

[Vulnerabilities By Type](#)

**Reports :**

[CVSS Score Report](#)

[CVSS Score Distribution](#)

**Search :**

[Vendor Search](#)

[Product Search](#)

[Version Search](#)

[Vulnerability Search](#)

[By Microsoft References](#)

## [Linux](#) » [Linux Kernel](#) : Security Vulnerabilities (CVSS Score between 7 and 7)

CVSS Scores Greater Than: 0 1 2 3 4 5 6 7 8 9

Sort Results By : [CVE Number Descending](#) [CVE Number Ascending](#) [CVSS Score](#) [Number of Exploits](#)

Total number of vulnerabilities : **268** Page : [1](#) (This Page) [2](#) [3](#) [4](#) [5](#)

[Copy Results](#) [Download Results](#) [Select Table](#)

#	CVE ID	CWE ID	# of Exploits	Vulnerability Type(s)	Published	Updated	Score
---	--------	--------	---------------	-----------------------	-----------	---------	-------

1	<a href="#">CVE-2015-5364</a>	<a href="#">399</a>		DoS	2015-08-31	2015-08-31	<b>7.8</b>
---	-------------------------------	---------------------	--	-----	------------	------------	------------

The (1) udp\_recvmmsg and (2) udpv6\_recvmmsg functions in the Linux kernel before 4.0.6 do not properly handle a packet flood via incorrect checksums within a UDP packet flood.

2	<a href="#">CVE-2015-5157</a>	<a href="#">264</a>		+Priv	2015-08-31	2015-08-31	<b>7.2</b>
---	-------------------------------	---------------------	--	-------	------------	------------	------------

arch/x86/entry/entry\_64.S in the Linux kernel before 4.1.6 on the x86\_64 platform mishandles IRET

					2015-08-31	2015-09-01	<b>7.2</b>
--	--	--	--	--	------------	------------	------------

drivers/vhost/scsi.c in the Linux kernel before 3.16.0 mishandles the VHOST\_SCSI\_SET\_ENDPOINT ioctl call. NOTE: This vulnerability affects only the Linux kernel.

					2015-06-07	2015-06-08	<b>7.8</b>
--	--	--	--	--	------------	------------	------------

zwpn/ozusbsvc1.c in the OZWPN driver in the Linux kernel before 3.16.0 does not properly handle a packet flood via incorrect checksums within a UDP packet flood.

					2015-08-31	2015-08-31	<b>7.2</b>
--	--	--	--	--	------------	------------	------------

6 on the x86\_64 platform improperly relies on the x86\_64 platform to properly handle a packet flood via incorrect checksums within a UDP packet flood.

					2015-08-08	2015-08-25	<b>7.2</b>
--	--	--	--	--	------------	------------	------------

is/pipe.c in the Linux kernel before 3.16.0 does not properly handle a packet flood via incorrect checksums within a UDP packet flood, which can cause a denial of service (system crash) or a denial of service (system hang).



# What can you do?

- A User



# What can you do?

- A technician





# What can you do?

- An administrator



# What can you do?

- A system programmer



# Also you can make the system



Minhyong Lee 06.29.2012



A close-up, low-angle shot of a tachometer. The needle is illuminated with a bright orange glow and points to the number 8 on the scale. The scale itself is also illuminated with orange light, with numbers 1 through 10 visible. The background is dark, making the glowing elements stand out. A blue rectangular box with white text is overlaid on the bottom half of the image.

Start your engine!

x1000r/min


# UNION

00011110 00011110 00011110 00011110 00011110



# Why talking about UNIX?


- Linux\* is a member of the large family of Unix-like operating systems



# Relationship between UNIX and Linux

## UNIX


Ken Thompson, Dennis Ritchie, Douglas  
McIlroy, and Joe Ossanna




## MINIX

Andrew S. Tanenbaum

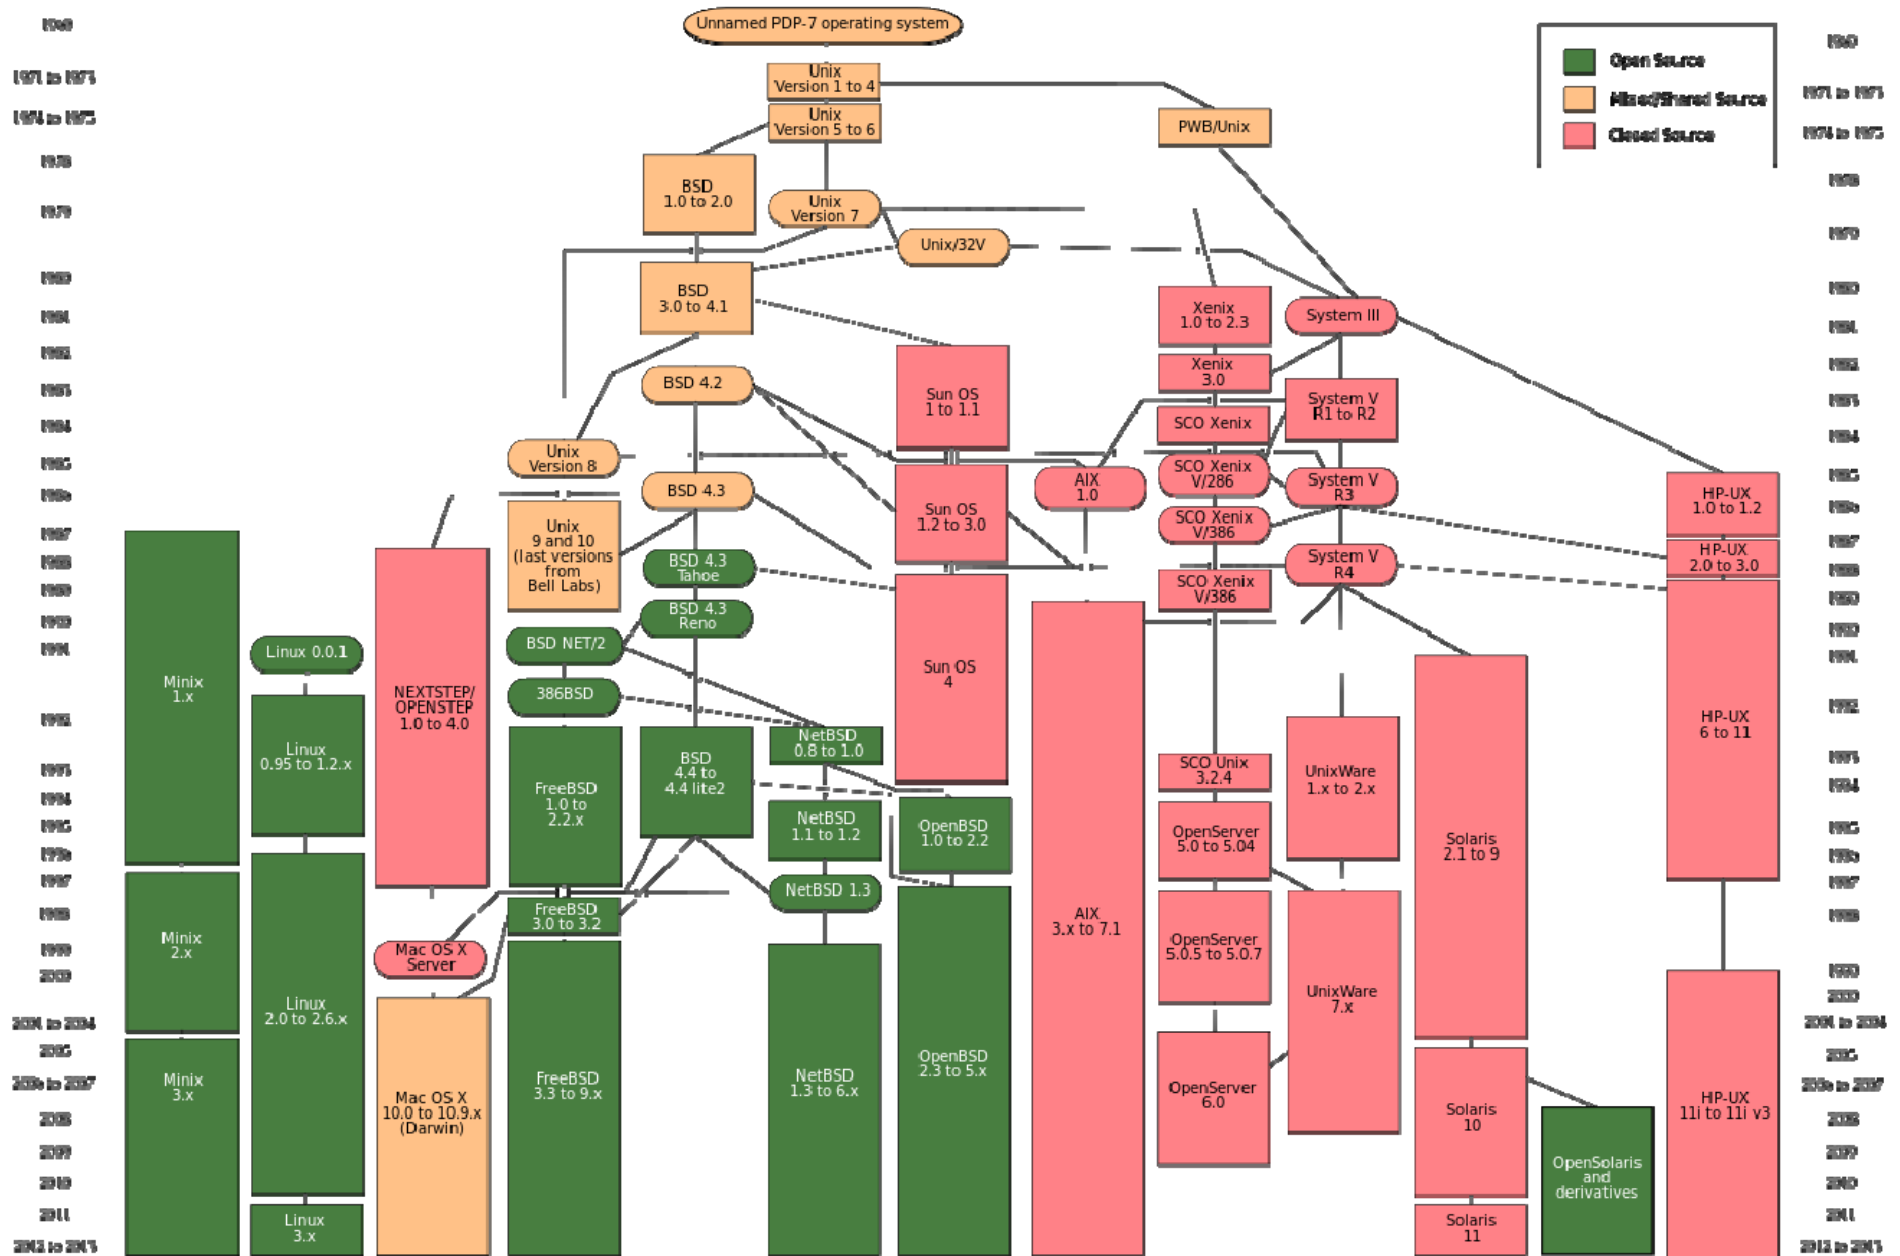
## Linux™



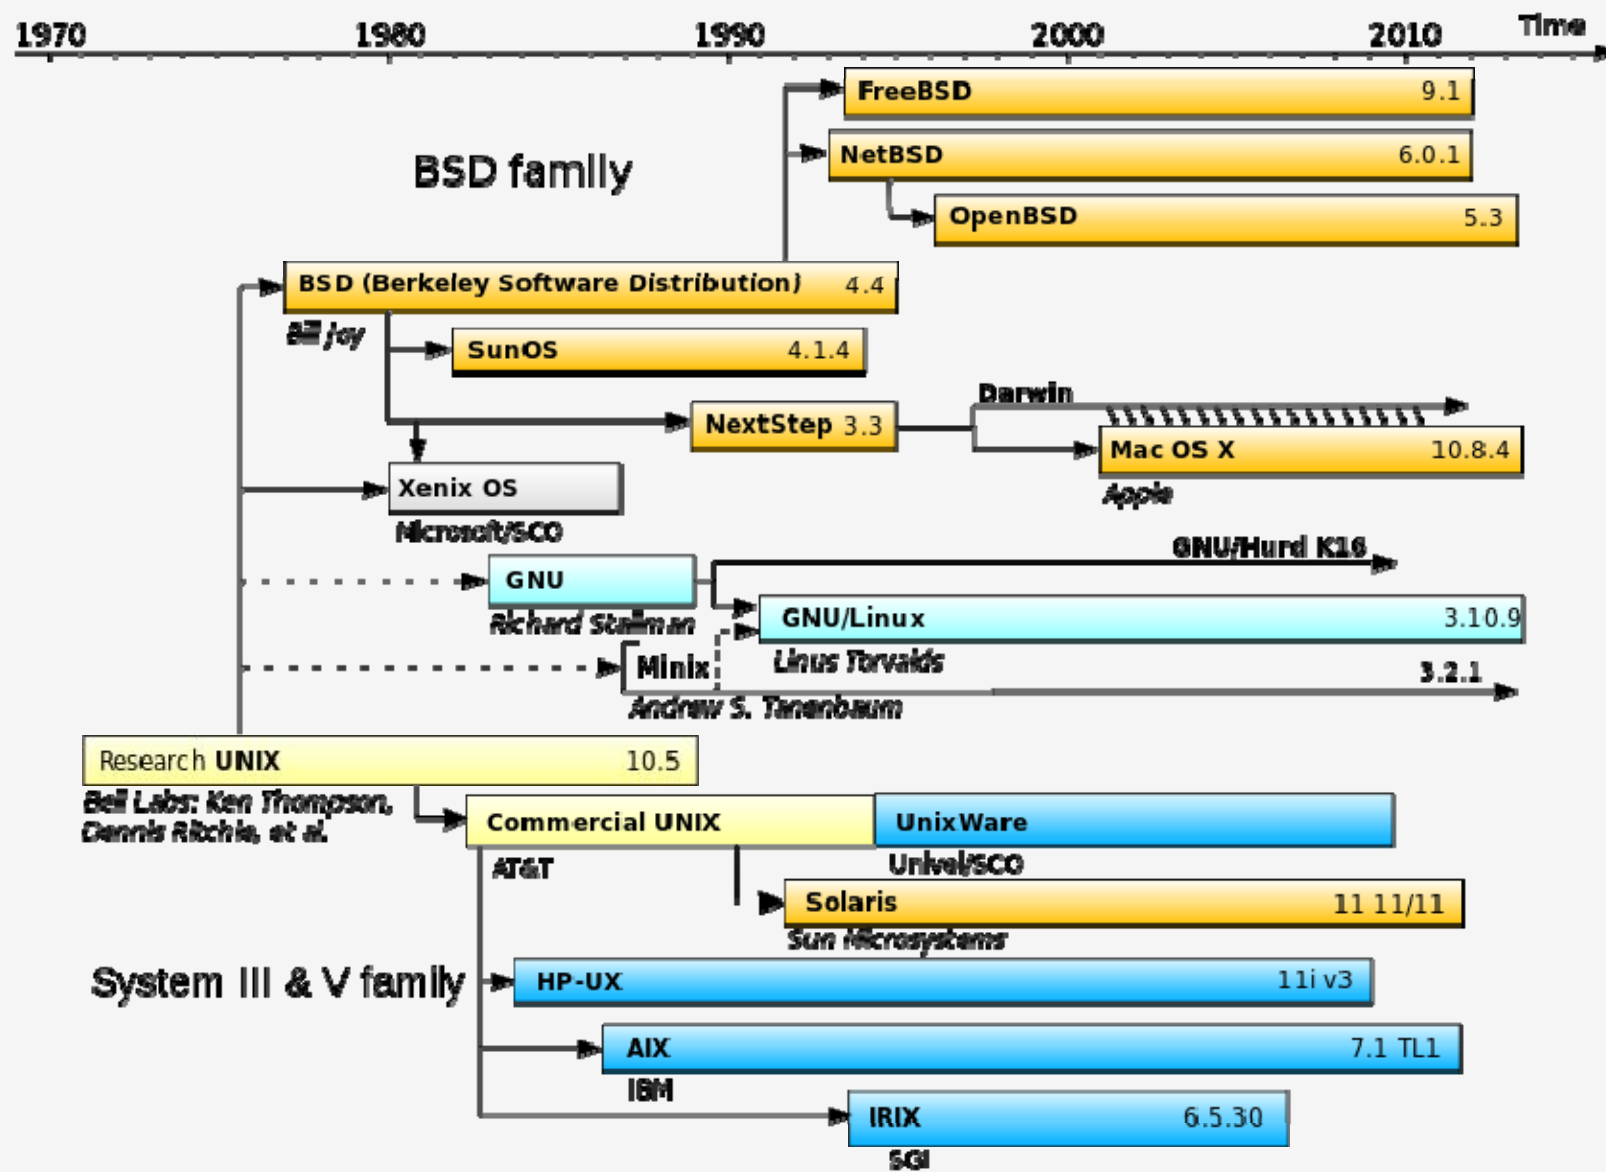
Linus Torvalds



# Unix Like OSes







# **What's common on the UNIX like OSes? – Similar philosophy**

**“Unix is simple and coherent, but it takes a genius (or at any rate a programmer) to understand and appreciate the simplicity.”**

**— Dennis Ritchie**

**“Note from the authors: Yes, we have lost our minds. Be forewarned: You will lose yours too.”**

**— Benny Goodheart & James Cox**

# What's common on the UNIX like OSes? – Similar philosophy

- Characterized by a **modular** design that is sometimes called the “**Unix philosophy**”:
  - the operating system provides **a set of simple tools** that each perform a limited, well-defined function
  - with a **unified filesystem** as the main means of **communication**
  - and a **shell scripting and command language** to **combine** the tools to perform complex workflows.

# What are differences?

## UNIX

Monolithic Kernel

## MINIX

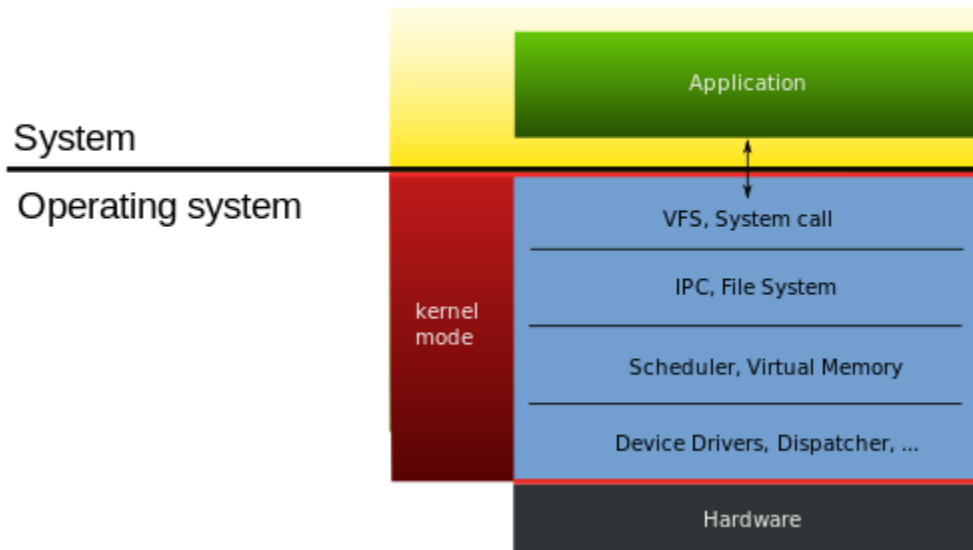
Microkernel

## Linux™



Monolithic Kernel (modular)

### Monolithic Kernel based Operating System



Examples: [UNIX/Linux](#)

Pro: Better performance

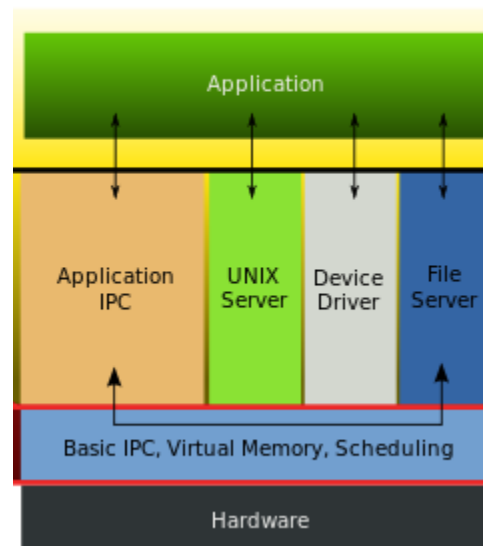
Con: Large and complex kernel

Example: [MINIX](#)

Pro: Demand  
very small set  
of functions  
Easily ported to  
another system

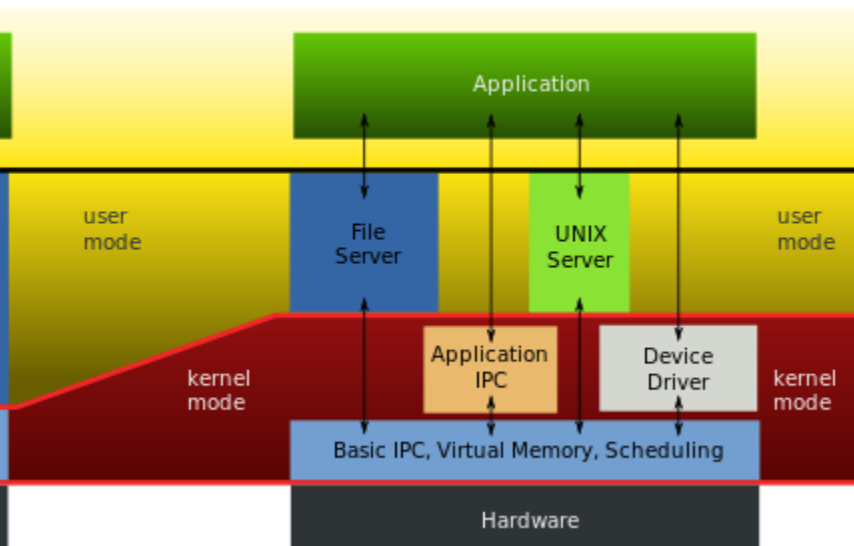
Con: Slower

### Microkernel based Operating System



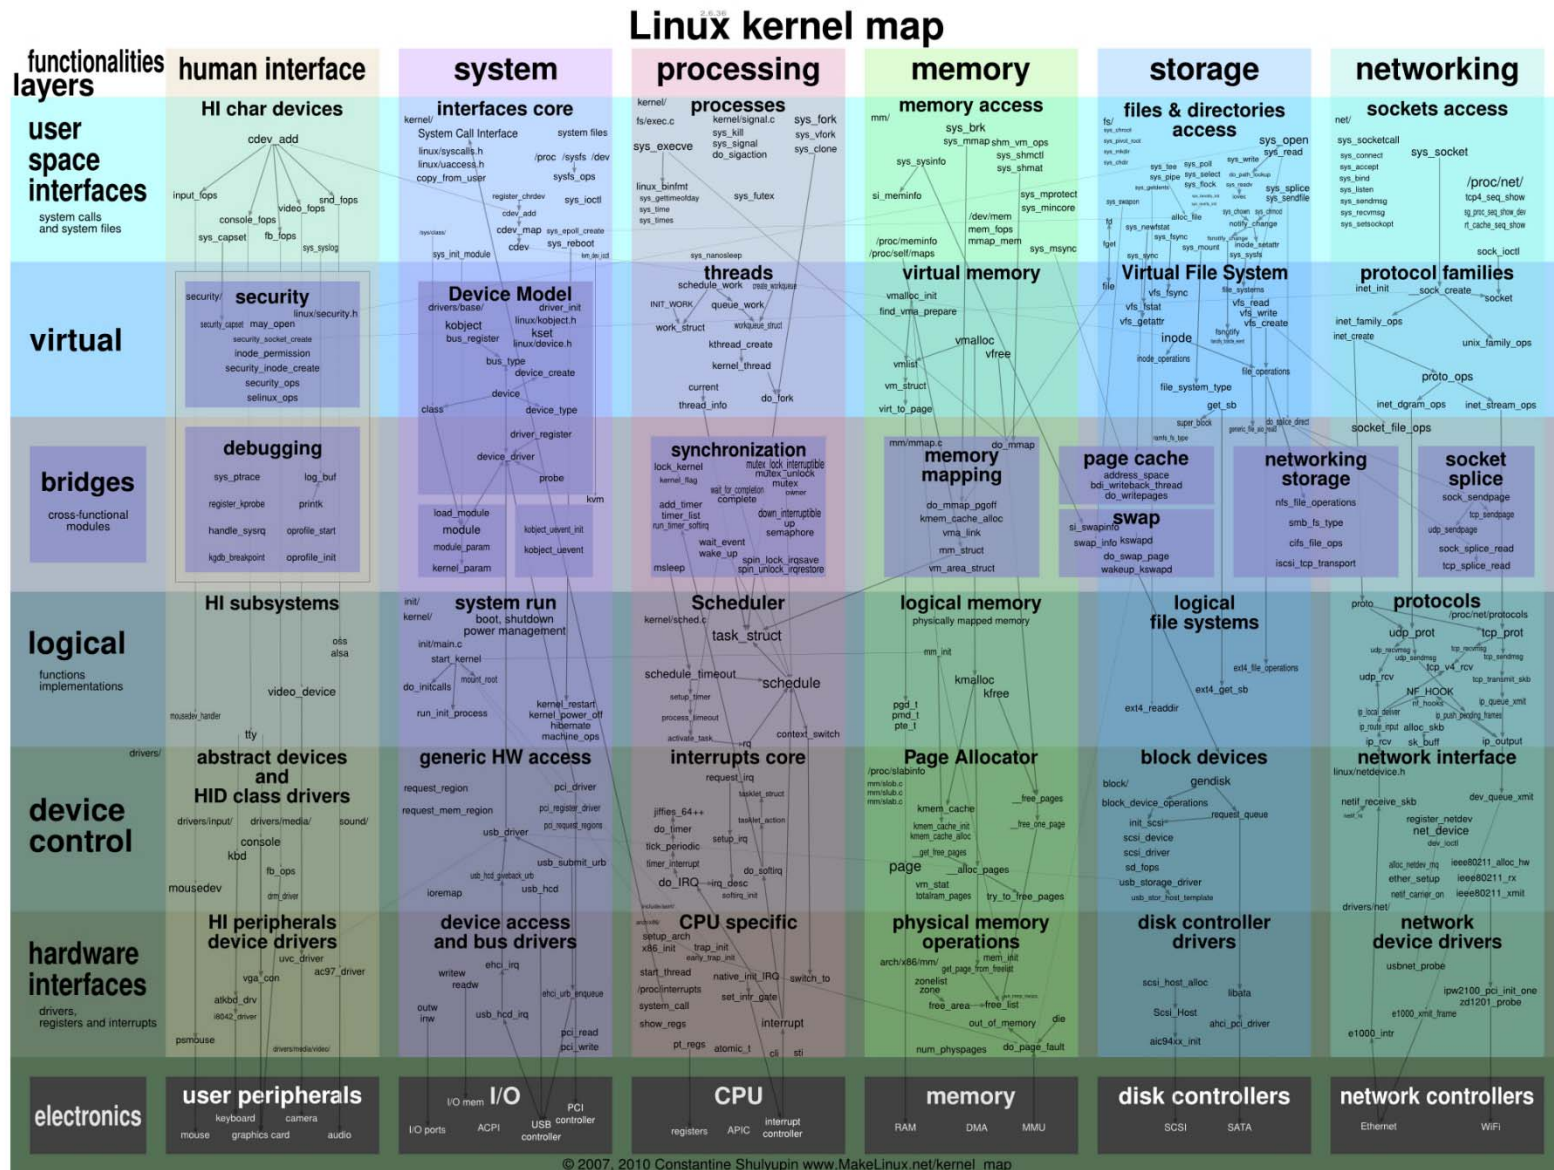
Example: [Windows](#)

### "Hybrid kernel" based Operating System

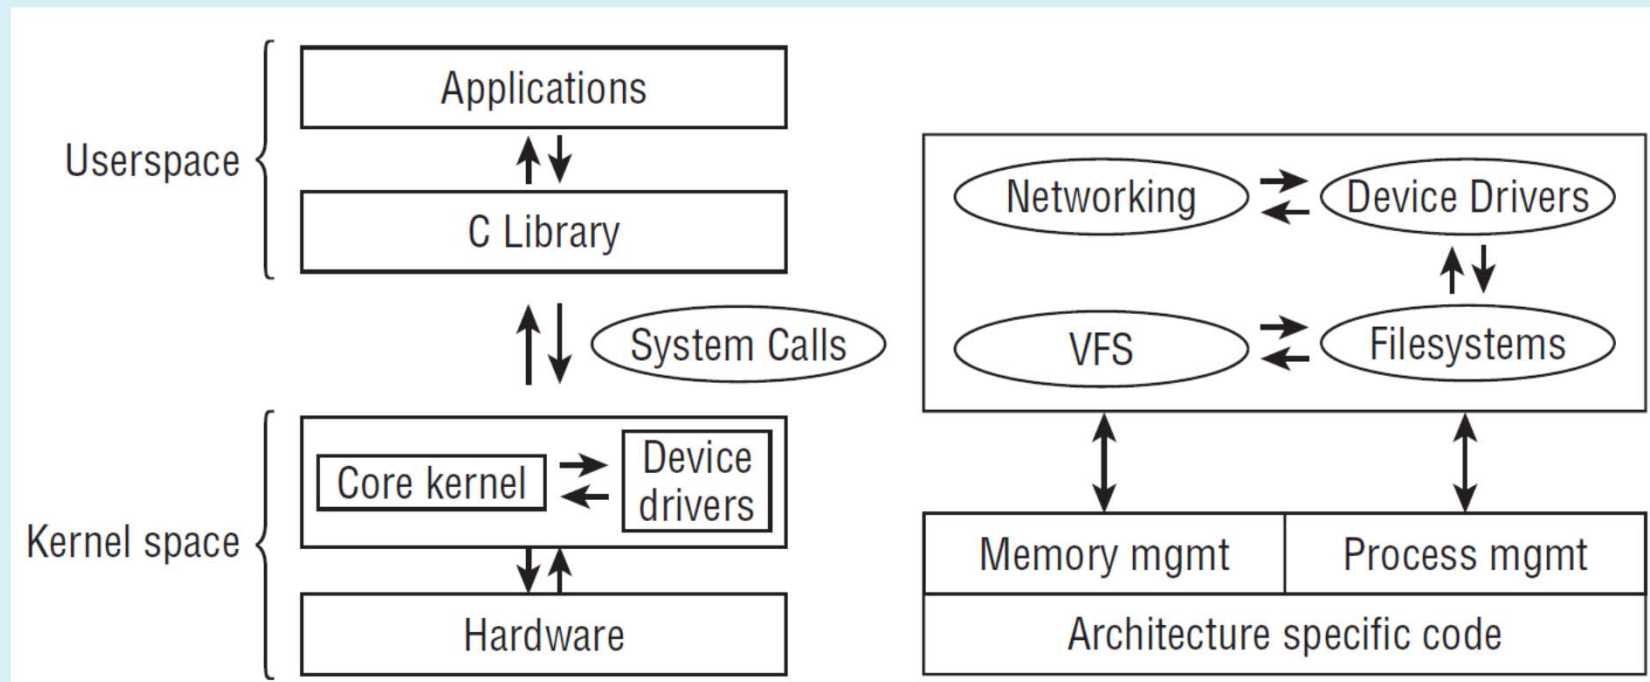




# Linux Kernel Architecture

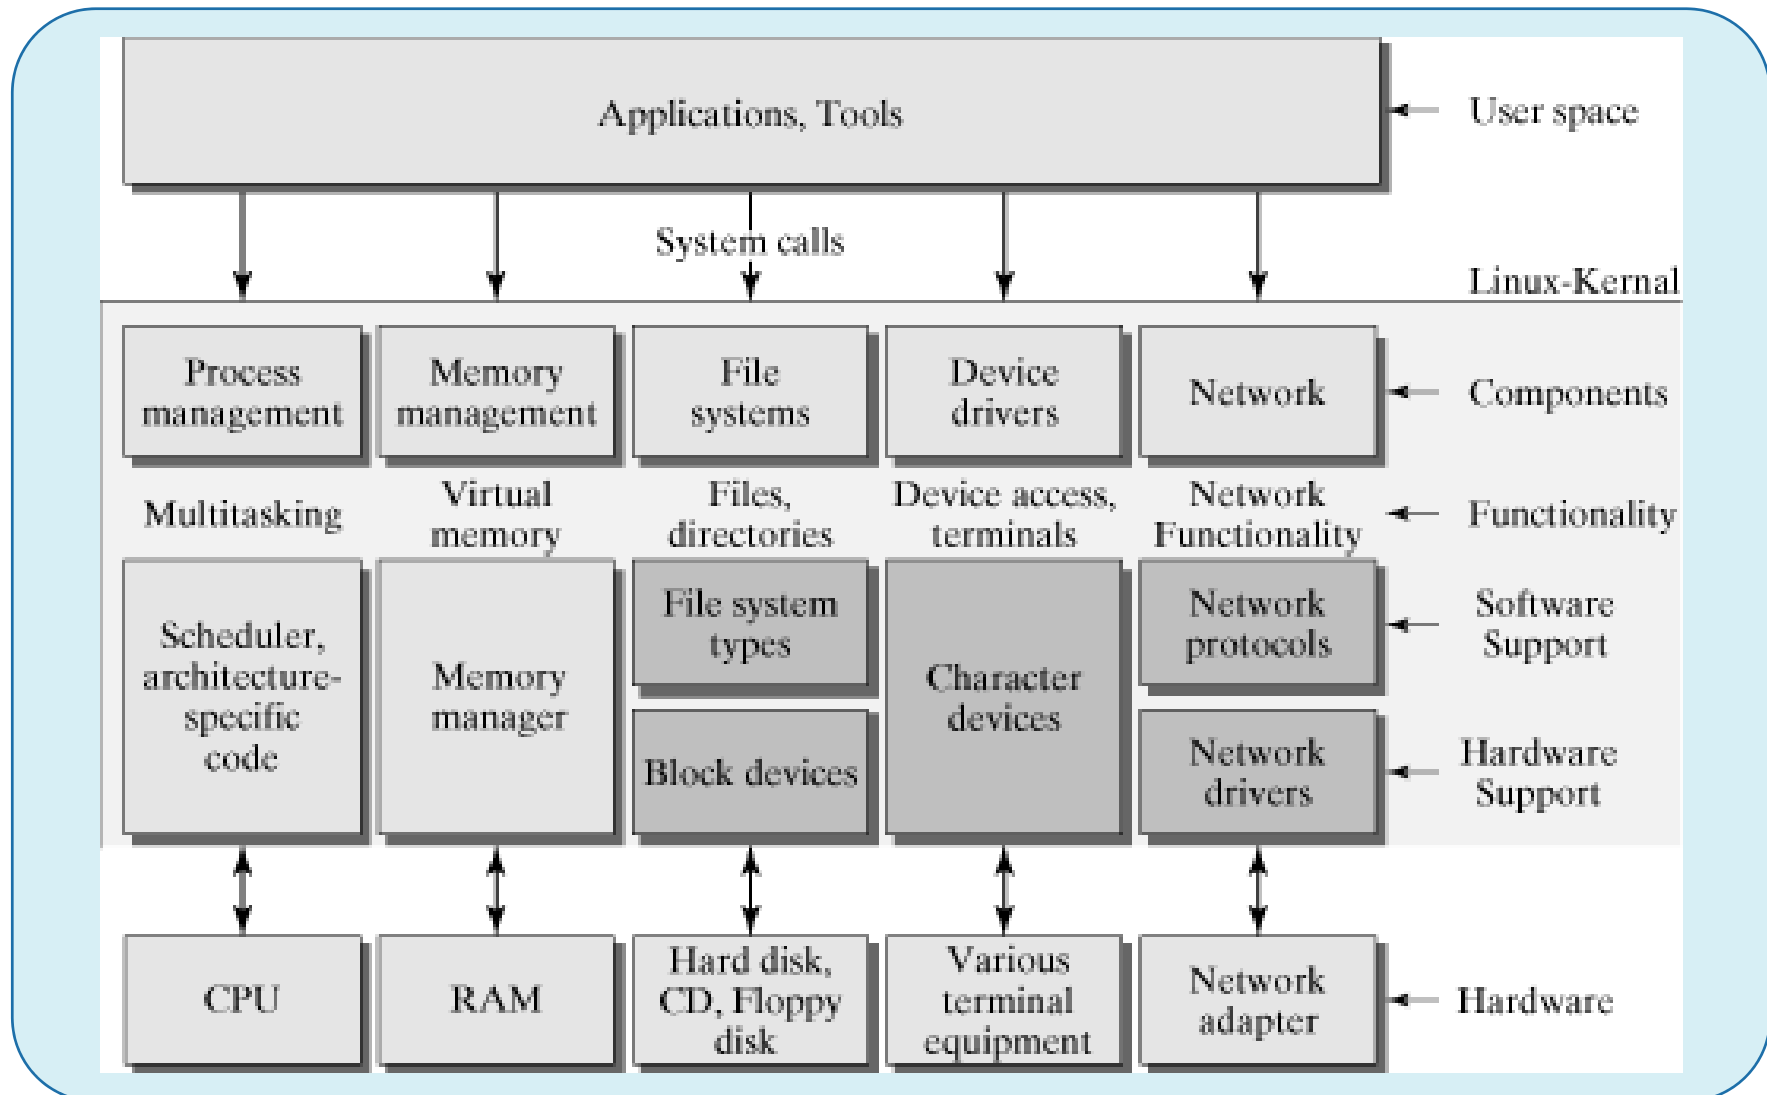


# Linux Kernel Architecture



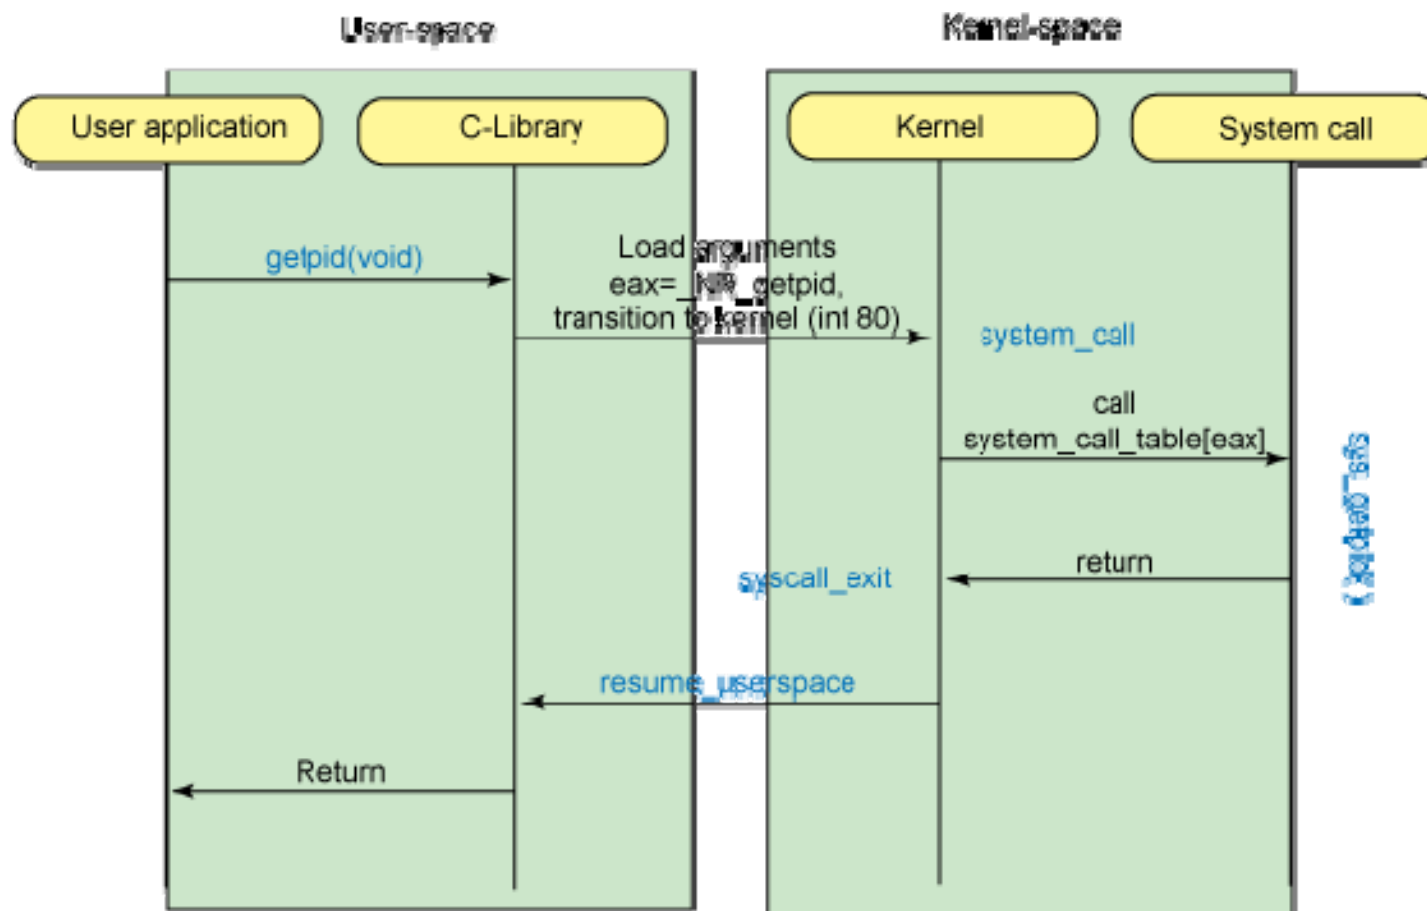
High-level view of Linux Kernel Architecture

# Linux Kernel Architecture



# What's going on?

- When application try to get its process id



# What's going on?

- Let's see in much lower level.

hello.asm

```
section .data
    msg db    "hello, world!"
```

```
section .text
    global _start
```


```
_start:
```

```
    mov     rax, 1
    mov     rdi, 1
    mov     rsi, msg
    mov     rdx, 13
    syscall

    mov     rax, 60
    mov     rdi, 0
    syscall
```

To assemble and run

```
$ nasm -f elf64 -o hello.o hello.asm
$ ld -o hello hello.o
```




%rax	System call	%rdi	%rsi	%rdx
0	sys_read	unsigned int fd	char *buf	size_t count
1	sys_write	unsigned int fd	const char *buf	size_t count

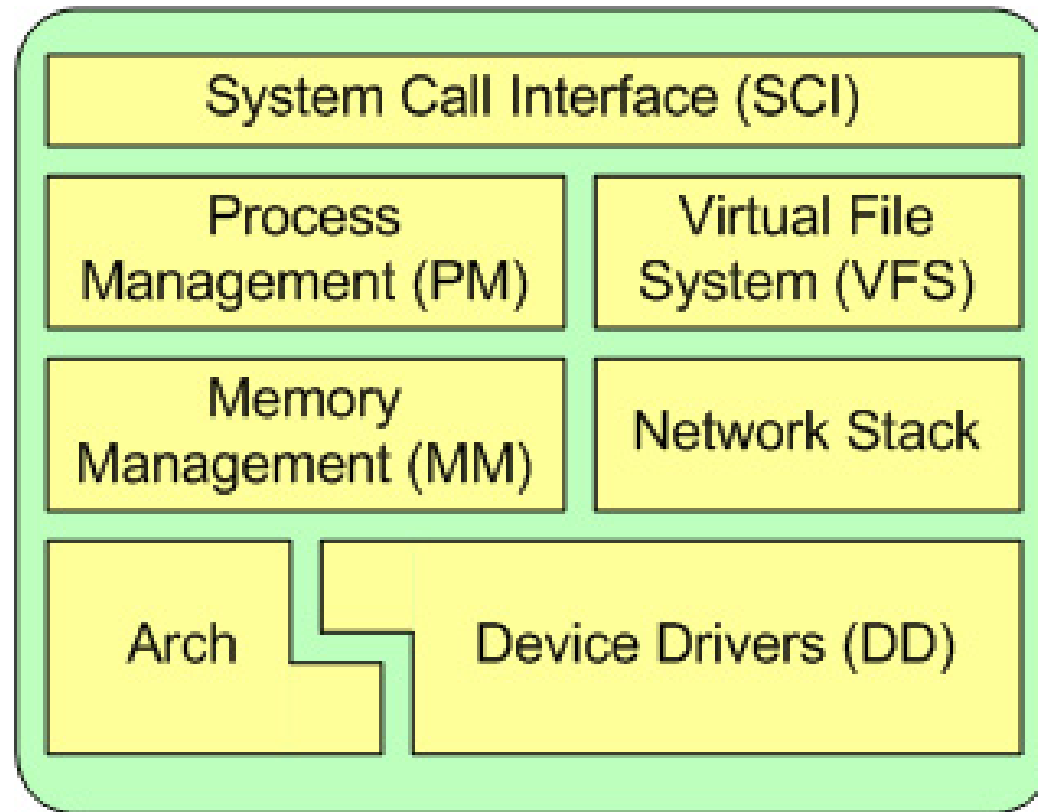
60	sys_exit	int error_code
----	----------	----------------

Linux system call table can be found:  
[http://blog.rchapman.org/post/36801038863/  
linux-system-call-table-for-x86-64](http://blog.rchapman.org/post/36801038863/linux-system-call-table-for-x86-64)



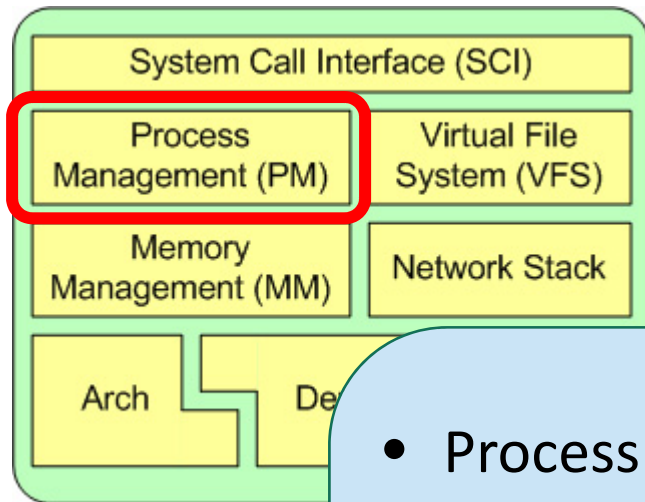
# OS Privilege Mode





**Let's take a look at the Linux kernel**

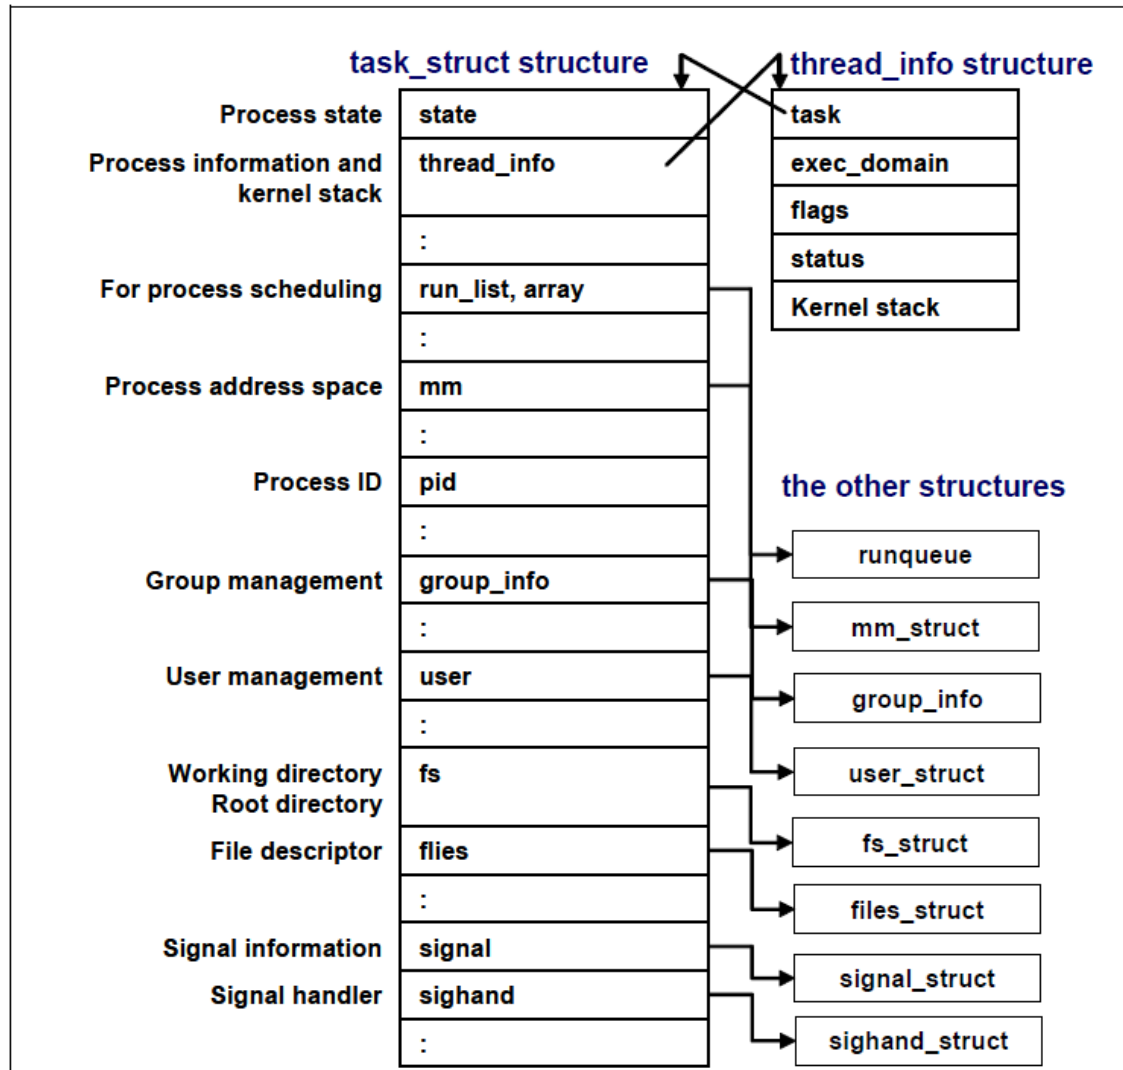
# Process Management




- Process management is focused on the execution of processes.
- In the kernel, these are called threads and represent an individual virtualization of the processor (thread code, data, stack, and CPU registers).

# Process Management


- Process information structure for Linux



# Process Management



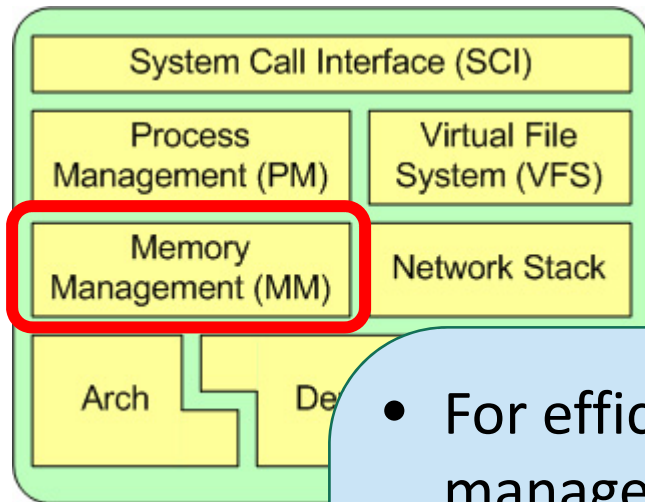
- Process states transitions



PROCESS




# Memory Management



- For efficiency, given the way that the hardware manages virtual memory, memory is managed in what are called pages (4KB in size for most architectures).
- Linux includes the means to manage the available memory, as well as the hardware mechanisms for physical and virtual mappings.

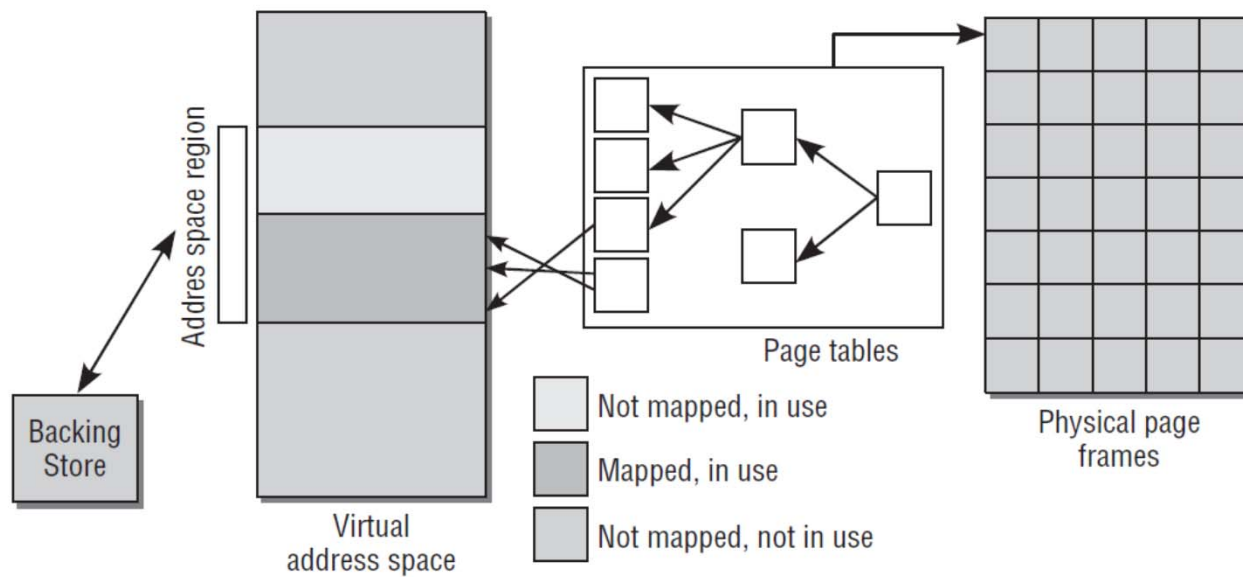
# Memory Management



So Kernel + User Spaces add for 256 TiB which is a tiny part of the 16 777 216 TiB addressable over 64 bit!

64 bit virtual memory address layout

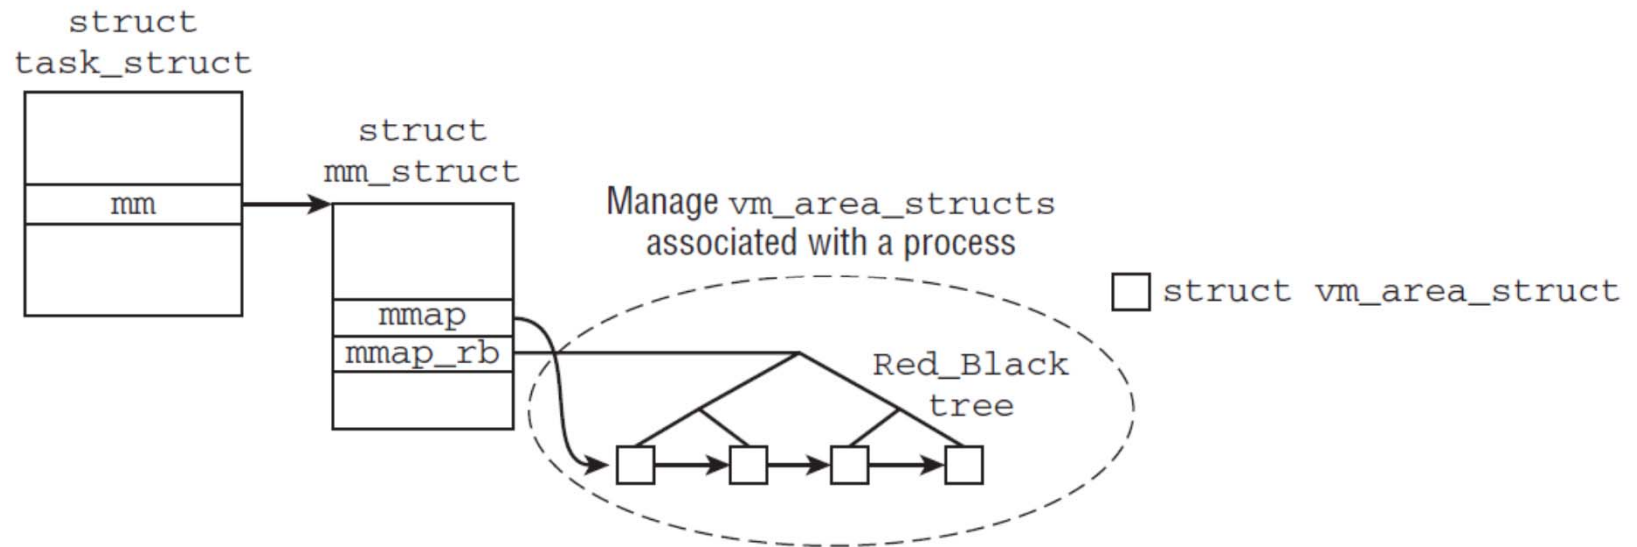
# Memory Management



Virtual memory paging




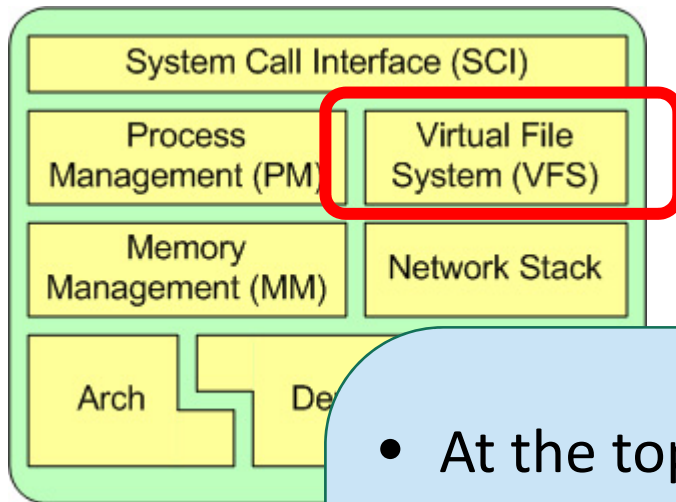
# Memory Management



- Virtual memory structures

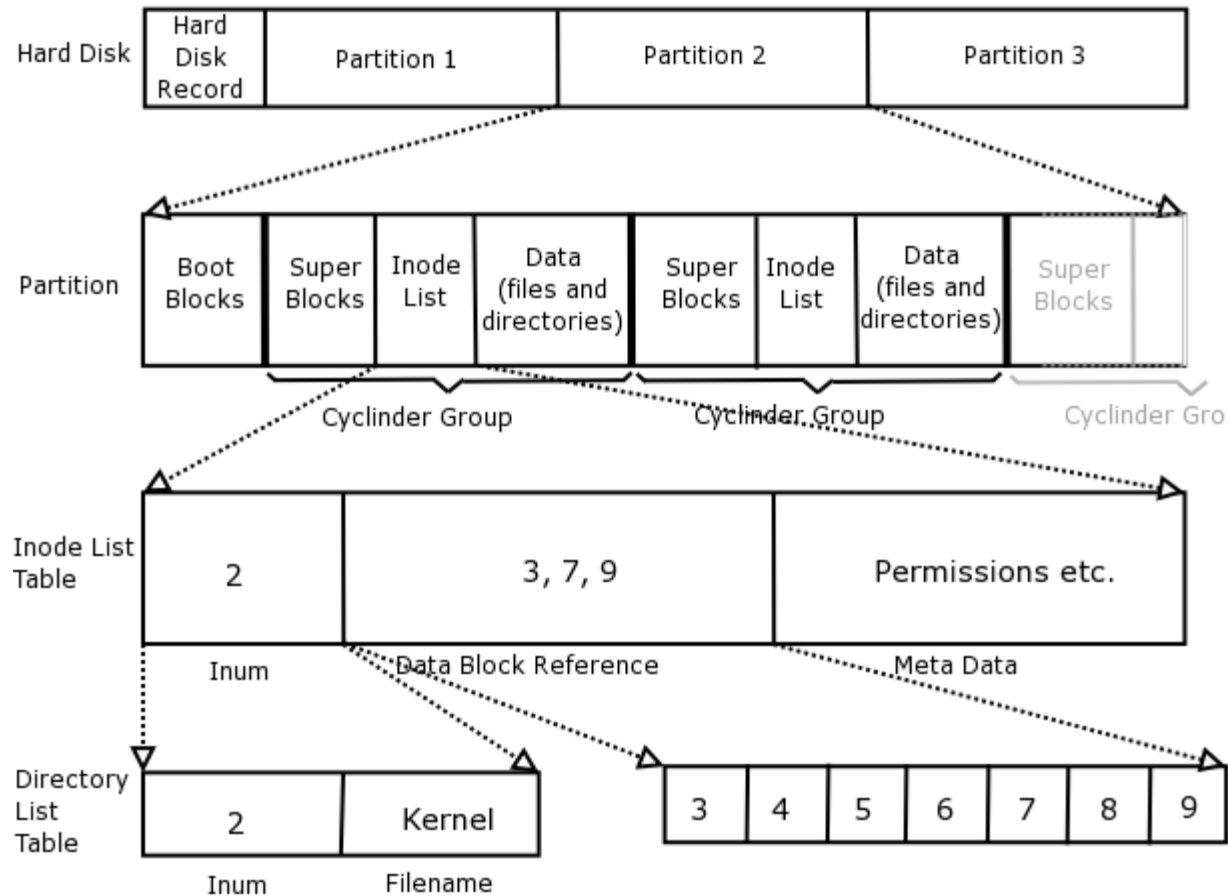


# File System



- At the top of the VFS is a common API abstraction of functions such as open, close, read, and write.
- At the bottom of the VFS are the file system abstractions that define how the upper-layer functions are implemented. These are plug-ins for the given file system (of which over 50 exist).

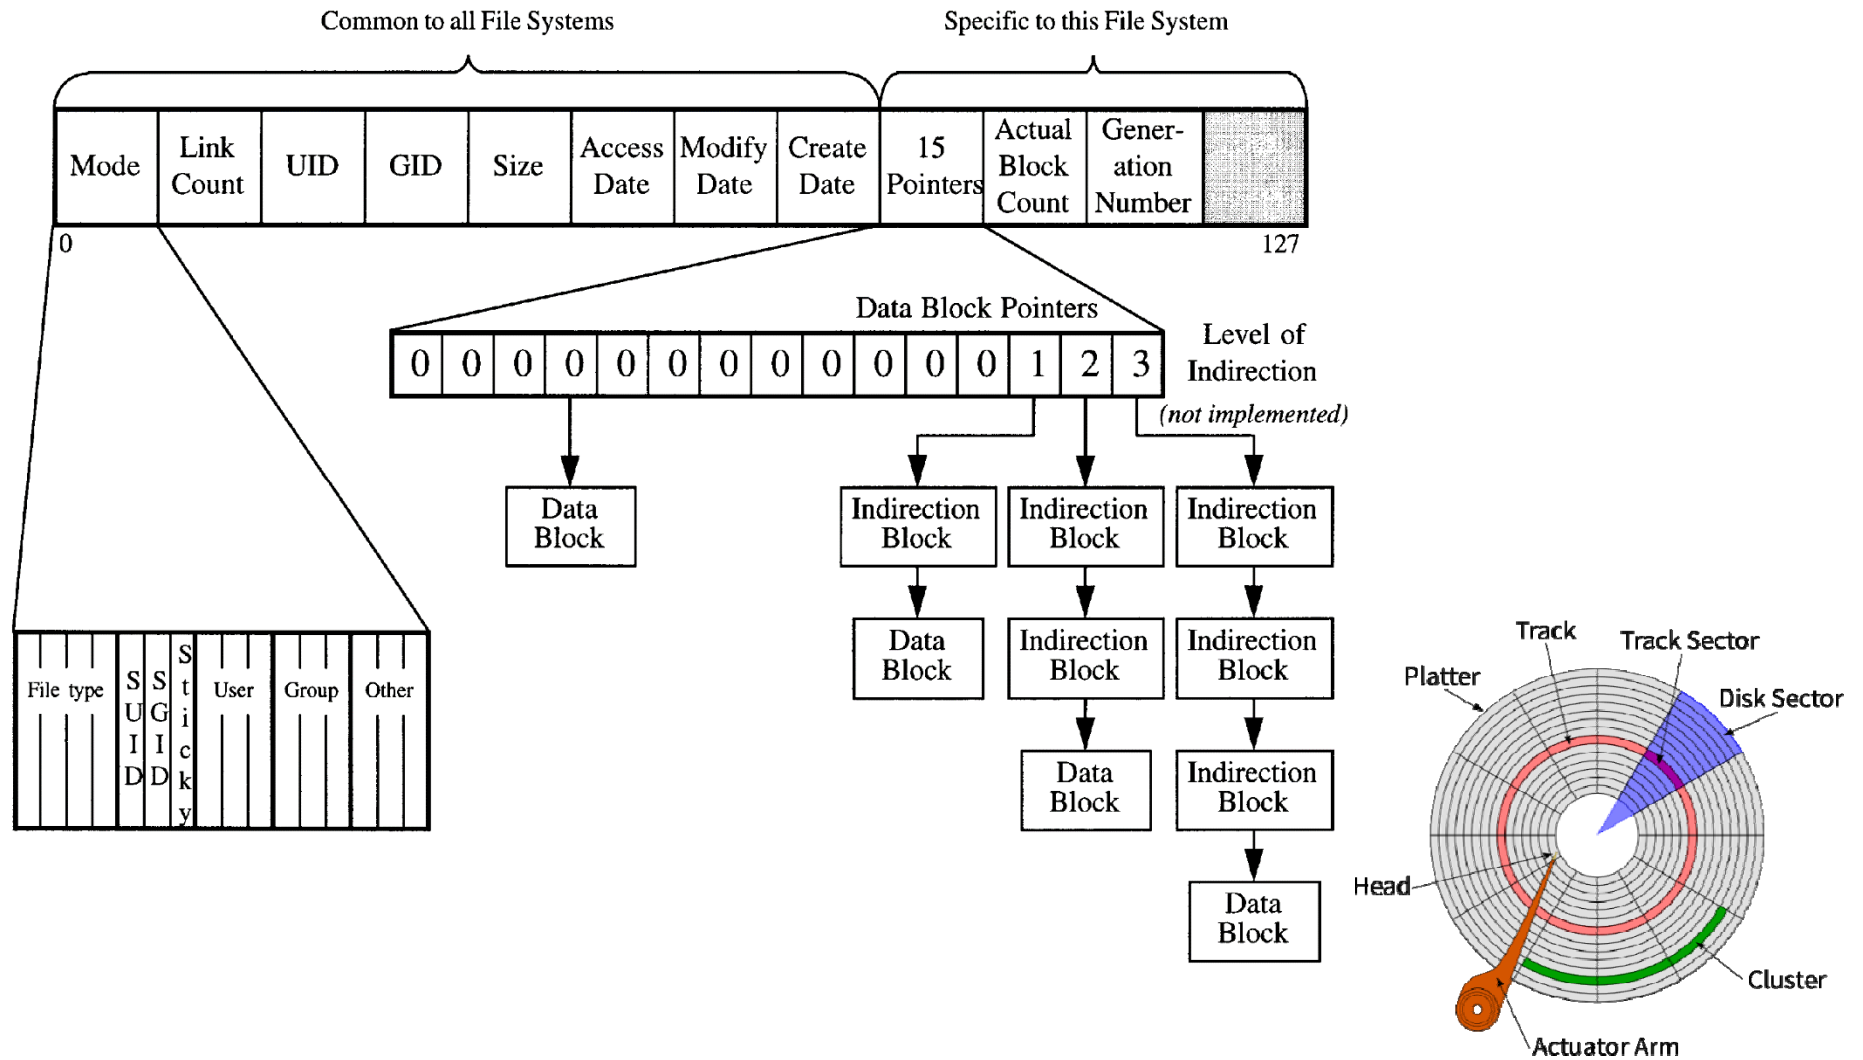
# File System Layout



File System Layout



# Disk Inode






# How this knowledge helps?

- Performance tuning
  - Change memory paging size
    - Recompile the kernel with a proper page size options
  - Change disk block size
    - Create a new file system with different options

Block Size	Cached Normal	Cached Direct	Uncached Normal	Uncached Direct
1	5.887	(n/a)	8306	(n/a)
27	6.135	(n/a)	8479	(n/a)
512	8.713	2498	8742	2579
2048	9.733	2422	8670	2721
4096	10.350	2514	25.039	2560

## Results:

no bs= 78s 144584+0 records  
bs=512 78s 144584+0 records  
bs=1k 38s 72292+0 records  
bs=2k 38s 36146+0 records  
bs=4k 38s 18073+0 records  
bs=5k 39s 14458+1 records  
bs=50k 38s 1445+1 records  
bs=500k 39s 144+1 records  
bs=512k 39s 144+1 records  
bs=1M 39s 72+1 records  
bs=5M 39s 14+1 records  
bs=10M 39s 7+1 records



# How this knowledge helps?

- Performance tuning
  - Use customizable Linux distro



# How this knowledge helps?

- Performance tuning
  - Unload unnecessary modules

## Module Commands

**depmod** - handle dependency descriptions for loadable kernel modules.

**insmod** - install loadable kernel module.

**lsmod** - list loaded modules.

**modinfo** - display information about a kernel module.

**modprobe** - high level handling of loadable modules.

**rmmod** - unload loadable modules.

# How this knowledge helps?

- Performance tuning
  - Unload unnecessary modules

```
$ lsmod
Module                Size  Used by
pci_stub              12622  1
joydev                17381  0
hid_generic           12548  0
usbhid                52659  0
hid                   106148  2 hid_generic,usbhid
eeepc_wmi              13151  0
parport_pc            32701  0
```

- **Secure Systems**



- **Secure Systems**

# How this knowledge helps?

- Kernel hardening
  - Disable and blacklist Linux modules



## Ghost in the Machine: Linux Zero-Day Vulnerability Opens Door for Attack

BY PAMELA COBB • JANUARY 29, 2015

# Questions?



Thank you

